

Wasiq Nabi Bakhsh

Buffalo, NY • +1 716-256-5828 • Wasiqcyber@gmail.com • LinkedIn • GitHub • Portfolio

Professional Summary

Data Engineer with 3+ years of experience designing scalable ETL/ELT pipelines, cloud data infrastructure, and BI platforms across banking, franchise analytics, and healthcare. MS Engineering Data Science candidate at University at Buffalo (Dec 2026). Key results: **1M+** transaction pipeline, **50%** latency reduction, **40%** CRM automation gain. Core stack: Python, SQL, PySpark, Databricks, Snowflake, Delta Lake, AWS, Tableau.

Technical Skills

Languages: Python, SQL, R, Bash, PostgreSQL, MySQL
Data Engineering: Apache Spark, PySpark, Delta Lake, dbt, ETL/ELT Pipelines, Airflow, Metadata-Driven Ingestion
Cloud and Storage: AWS S3/Lambda, GCP, Databricks, Snowflake, Azure DevOps, CI/CD
BI and Visualization: Tableau, Power BI, Salesforce (Einstein AI, Flow Builder), Looker Studio, Matplotlib, Seaborn
ML and MLOps: Scikit-learn, TensorFlow, XGBoost, MLflow, Docker, FastAPI, LangGraph, RAG Pipelines
Certifications: Google Data Analytics | AWS Data Analytics | Tableau Visualization (Honors) | IBM Data Analysis with Python

Education

University at Buffalo – State University of New York Buffalo, NY
MS in Engineering Data Science Aug 2025 – Dec 2026

Experience

University at Buffalo – International Student Services Buffalo, NY
Data Science Graduate Assistant Jan 2026 – Present

- **Built ingestion and validation pipelines** (Python/SQL) for **500+** student records on AWS S3/Lambda, enforcing schema checks, type casting, deduplication, and null-threshold rules, enabling faster data onboarding and reduced data-entry errors.
- **Designed Tableau dashboards** tracking retention and service KPIs; standardized data preparation workflows reducing manual reporting effort by **30%**.
- Partnered with staff to surface cohort-level insights, **translating pipeline outputs into actionable student support strategies**.

Franchise Pipeline Solutions Chicago, IL (Remote)
Data Engineer Jan 2024 – Jun 2025

- **Built lead-to-deal data pipelines** in Salesforce processing incoming account data and engineered ETL workflows integrating Salesforce CRM with Tableau to consolidate multi-source data into a unified analytics layer for reporting and forecasting.
- **Developed revenue forecasts, budget projections and agent performance dashboards** using time-series modeling in Python and Salesforce reporting measuring conversion rates, deal velocity and revenue contribution per agent.
- **Automated Excel-based reporting workflows using AI and Python** eliminating manual processes and reduced reporting time by **40%** and delivered executive presentations translating pipeline analytics into business recommendations.

The Bank of Khyber Pakistan
Data Science Intern Jul 2024 – Jul 2025

- **Built reproducible ETL pipelines** across SQL and AWS S3 ingesting and transforming **1M+** transactions, applying feature engineering (merchant risk scores, velocity flags, amount deviations) for fraud analytics.
- **Developed fraud detection models** (XGBoost, Neural Networks) improving accuracy by **30%**; automated dashboards cutting fraud response time by **15%**.

House of Brands Media Remote
Data Scientist Feb 2024 – Oct 2024

- **Automated data preparation pipelines** reducing operational latency by **50%**; built **A/B testing and lift analysis frameworks** improving lead generation by **25%** and driving **30%** revenue growth.
- Designed **time-series models with Apache Spark** to evaluate marketing channel mix; delivered **Snowflake dataset integrations** consolidating multi-source campaign data for downstream analytics.

Projects

Lakehouse Data Engineering Pipeline | *Databricks, Delta Lake, PySpark, Snowflake, Azure DevOps* 2025

- **Enterprise-grade lakehouse on Databricks** with metadata-driven ingestion, **PySpark transformations**, and **Kimball Star Schema warehouse on Snowflake**; CI/CD automated via Azure DevOps.

End-to-End ML Deployment Pipeline | *Python, FastAPI, Docker, MLflow, DagsHub, Render* 2025

- **Full data-to-deployment pipeline** with **16** experiments tracked via MLflow/DagsHub; best model served via **FastAPI and Streamlit**, Dockerized and deployed to Render cloud.

Financial Sentiment Data Pipeline | *FinBERT, XGBoost, LSTM, Python, NLP* 2025

- **Scalable pipeline** aligning **600K+** financial news headlines with **7,000+** U.S. stocks; engineered **temporal features** (rolling averages, momentum) and applied SHAP for model interpretability.